

Optional Lab - Regularized Cost and Gradient

Goals ¶

In this lab, you will:

- extend the previous linear and logistic cost functions with a regularization term.
- rerun the previous example of over-fitting with a regularization term added.

In []:

```
import numpy as np
%matplotlib widget
import matplotlib.pyplot as plt
from plt_overfit import overfit_example, output
from lab_utils_common import sigmoid
np.set_printoptions(precision=8)
```

Adding regularization



The slides above show the cost and gradient functions for both linear and logistic regression. Note:

- Cost
 - The cost functions differ significantly between linear and logistic regression, but adding regularization to the equations is the same.
- Gradient
 - The gradient functions for linear and logistic regression are very similar. They differ only in the implementation of f_{wb} .

Cost functions with regularization

Cost function for regularized linear regression

The equation for the cost function regularized linear regression is:

$$J(\mathbf{w}, b) = \frac{1}{2m} \sum_{i=0}^{m-1} (f_{\mathbf{w},b}(\mathbf{x}^{(i)}) - y^{(i)})^2 + \frac{\lambda}{2m} \sum_{j=0}^{n-1} w_j^2 \quad (1)$$

where:

$$f_{\mathbf{w},b}(\mathbf{x}^{(i)}) = \mathbf{w} \cdot \mathbf{x}^{(i)} + b \quad (2)$$

Compare this to the cost function without regularization (which you implemented in a previous lab), which is of the form:

$$J(\mathbf{w}, b) = \frac{1}{2m} \sum_{i=0}^{m-1} (f_{\mathbf{w},b}(\mathbf{x}^{(i)}) - y^{(i)})^2$$

The difference is the regularization term, $\frac{\lambda}{2m} \sum_{j=0}^{n-1} w_j^2$

Including this term encourages gradient descent to minimize the size of the parameters. Note, in this example, the parameter b is not regularized. This is standard practice.

Below is an implementation of equations (1) and (2). Note that this uses a *standard pattern for this course*, a `for` loop over all `m` examples.

In []:

```
def compute_cost_linear_reg(X, y, w, b, lambda_ = 1):
    """
    Computes the cost over all examples
    Args:
        X (ndarray (m,n)): Data, m examples with n features
        y (ndarray (m,)): target values
        w (ndarray (n,)): model parameters
        b (scalar)       : model parameter
        lambda_ (scalar): Controls amount of regularization
    Returns:
        total_cost (scalar): cost
    """

    m = X.shape[0]
    n = len(w)
    cost = 0.
    for i in range(m):
        f_wb_i = np.dot(X[i], w) + b                #(n,)(n,)
        =scalar, see np.dot
        cost = cost + (f_wb_i - y[i])**2            #scalar
    cost = cost / (2 * m)                           #scalar

    reg_cost = 0
    for j in range(n):
        reg_cost += (w[j]**2)                        #scalar
    reg_cost = (lambda_/(2*m)) * reg_cost            #scalar

    total_cost = cost + reg_cost                     #scalar
    return total_cost                               #scalar
```

Run the cell below to see it in action.

In []:

```
np.random.seed(1)
X_tmp = np.random.rand(5,6)
y_tmp = np.array([0,1,0,1,0])
w_tmp = np.random.rand(X_tmp.shape[1]).reshape(-1,)-0.5
b_tmp = 0.5
lambda_tmp = 0.7
cost_tmp = compute_cost_linear_reg(X_tmp, y_tmp, w_tmp, b_tmp, lambda_tmp)

print("Regularized cost:", cost_tmp)
```

Expected Output:

Regularized cost: 0.07917239320214275

Cost function for regularized logistic regression

For regularized **logistic** regression, the cost function is of the form

$$J(\mathbf{w}, b) = \frac{1}{m} \sum_{i=0}^{m-1} [-y^{(i)} \log(f_{\mathbf{w},b}(\mathbf{x}^{(i)})) - (1 - y^{(i)}) \log(1 - f_{\mathbf{w},b}(\mathbf{x}^{(i)}))] + \frac{\lambda}{2m} \sum_{j=0}^{n-1} w_j^2 \quad (3)$$

where:

$$f_{\mathbf{w},b}(\mathbf{x}^{(i)}) = \text{sigmoid}(\mathbf{w} \cdot \mathbf{x}^{(i)} + b) \quad (4)$$

Compare this to the cost function without regularization (which you implemented in a previous lab):

$$J(\mathbf{w}, b) = \frac{1}{m} \sum_{i=0}^{m-1} [(-y^{(i)} \log(f_{\mathbf{w},b}(\mathbf{x}^{(i)})) - (1 - y^{(i)}) \log(1 - f_{\mathbf{w},b}(\mathbf{x}^{(i)})))]$$

As was the case in linear regression above, the difference is the regularization term, which is $\frac{\lambda}{2m} \sum_{j=0}^{n-1} w_j^2$

Including this term encourages gradient descent to minimize the size of the parameters. Note, in this example, the parameter b is not regularized. This is standard practice.

In []:

```
def compute_cost_logistic_reg(X, y, w, b, lambda_ = 1):
    """
    Computes the cost over all examples
    Args:
    Args:
        X (ndarray (m,n)): Data, m examples with n features
        y (ndarray (m,)): target values
        w (ndarray (n,)): model parameters
        b (scalar)       : model parameter
        lambda_ (scalar): Controls amount of regularization
    Returns:
        total_cost (scalar): cost
    """

    m,n = X.shape
    cost = 0.
    for i in range(m):
        z_i = np.dot(X[i], w) + b                                #(n,)(n,)
        #scalar
        f_wb_i = sigmoid(z_i)                                    #scalar
        cost += -y[i]*np.log(f_wb_i) - (1-y[i])*np.log(1-f_wb_i) #scalar

    cost = cost/m                                                #scalar

    reg_cost = 0
    for j in range(n):
        reg_cost += (w[j]**2)                                    #scalar
    reg_cost = (lambda_/(2*m)) * reg_cost                        #scalar

    total_cost = cost + reg_cost                                  #scalar
    return total_cost                                            #scalar
```

Run the cell below to see it in action.

In []:

```
np.random.seed(1)
X_tmp = np.random.rand(5,6)
y_tmp = np.array([0,1,0,1,0])
w_tmp = np.random.rand(X_tmp.shape[1]).reshape(-1,)-0.5
b_tmp = 0.5
lambda_tmp = 0.7
cost_tmp = compute_cost_logistic_reg(X_tmp, y_tmp, w_tmp, b_tmp, lambda_tmp)

print("Regularized cost:", cost_tmp)
```

Expected Output:

Regularized cost: 0.6850849138741673

Gradient descent with regularization

The basic algorithm for running gradient descent does not change with regularization, it is:

repeat until convergence: {

$$w_j = w_j - \alpha \frac{\partial J(\mathbf{w}, b)}{\partial w_j} \quad \text{for } j := 0..n-1 \quad (1)$$

$$b = b - \alpha \frac{\partial J(\mathbf{w}, b)}{\partial b}$$

}

Where each iteration performs simultaneous updates on w_j for all j .

What changes with regularization is computing the gradients.

Computing the Gradient with regularization (both linear/logistic)

The gradient calculation for both linear and logistic regression are nearly identical, differing only in computation of $f_{\mathbf{w},b}$.

$$\frac{\partial J(\mathbf{w}, b)}{\partial w_j} = \frac{1}{m} \sum_{i=0}^{m-1} (f_{\mathbf{w},b}(\mathbf{x}^{(i)}) - y^{(i)})x_j^{(i)} + \frac{\lambda}{m}w_j \quad (2)$$

$$\frac{\partial J(\mathbf{w}, b)}{\partial b} = \frac{1}{m} \sum_{i=0}^{m-1} (f_{\mathbf{w},b}(\mathbf{x}^{(i)}) - y^{(i)}) \quad (3)$$

- m is the number of training examples in the data set
- $f_{\mathbf{w},b}(x^{(i)})$ is the model's prediction, while $y^{(i)}$ is the target

- For a **linear** regression model

$$f_{\mathbf{w},b}(x) = \mathbf{w} \cdot \mathbf{x} + b$$

- For a **logistic** regression model

$$z = \mathbf{w} \cdot \mathbf{x} + b$$

$$f_{\mathbf{w},b}(x) = g(z)$$

where $g(z)$ is the sigmoid function:

$$g(z) = \frac{1}{1+e^{-z}}$$

The term which adds regularization is the $\frac{\lambda}{m}w_j$.

Gradient function for regularized linear regression

In []:

```
def compute_gradient_linear_reg(X, y, w, b, lambda_):  
    """  
    Computes the gradient for linear regression  
    Args:  
        X (ndarray (m,n)): Data, m examples with n features  
        y (ndarray (m,)): target values  
        w (ndarray (n,)): model parameters  
        b (scalar)       : model parameter  
        lambda_ (scalar): Controls amount of regularization  
  
    Returns:  
        dj_dw (ndarray (n,)): The gradient of the cost w.r.t. the parameters w.  
        dj_db (scalar):       The gradient of the cost w.r.t. the parameter b.  
    """  
    m,n = X.shape                #(number of examples, number of features)  
    dj_dw = np.zeros((n,))  
    dj_db = 0.  
  
    for i in range(m):  
        err = (np.dot(X[i], w) + b) - y[i]  
        for j in range(n):  
            dj_dw[j] = dj_dw[j] + err * X[i, j]  
        dj_db = dj_db + err  
    dj_dw = dj_dw / m  
    dj_db = dj_db / m  
  
    for j in range(n):  
        dj_dw[j] = dj_dw[j] + (lambda_/m) * w[j]  
  
    return dj_db, dj_dw
```

Run the cell below to see it in action.

In []:

```
np.random.seed(1)  
X_tmp = np.random.rand(5,3)  
y_tmp = np.array([0,1,0,1,0])  
w_tmp = np.random.rand(X_tmp.shape[1])  
b_tmp = 0.5  
lambda_tmp = 0.7  
dj_db_tmp, dj_dw_tmp = compute_gradient_linear_reg(X_tmp, y_tmp, w_tmp, b_tmp,  
lambda_tmp)  
  
print(f"dj_db: {dj_db_tmp}", )  
print(f"Regularized dj_dw:\n {dj_dw_tmp.tolist()}", )
```

Expected Output

```
dj_db: 0.6648774569425726  
Regularized dj_dw:  
[0.29653214748822276, 0.4911679625918033, 0.21645877535865857]
```

Gradient function for regularized logistic regression

In []:

```
def compute_gradient_logistic_reg(X, y, w, b, lambda_):
    """
    Computes the gradient for linear regression

    Args:
        X (ndarray (m,n)): Data, m examples with n features
        y (ndarray (m,)): target values
        w (ndarray (n,)): model parameters
        b (scalar)       : model parameter
        lambda_ (scalar): Controls amount of regularization
    Returns
        dj_dw (ndarray Shape (n,)): The gradient of the cost w.r.t. the parameters
    w.
        dj_db (scalar)             : The gradient of the cost w.r.t. the parameter
    b.
    """
    m,n = X.shape
    dj_dw = np.zeros((n,))          #(n,)
    dj_db = 0.0                    #scalar

    for i in range(m):
        f_wb_i = sigmoid(np.dot(X[i],w) + b)          #(n,)(n,)=scalar
        err_i = f_wb_i - y[i]                         #scalar
        for j in range(n):
            dj_dw[j] = dj_dw[j] + err_i * X[i,j]      #scalar
        dj_db = dj_db + err_i
    dj_dw = dj_dw/m                      #(n,)
    dj_db = dj_db/m                      #scalar

    for j in range(n):
        dj_dw[j] = dj_dw[j] + (lambda_/m) * w[j]

    return dj_db, dj_dw
```

Run the cell below to see it in action.

In []:

```
np.random.seed(1)
X_tmp = np.random.rand(5,3)
y_tmp = np.array([0,1,0,1,0])
w_tmp = np.random.rand(X_tmp.shape[1])
b_tmp = 0.5
lambda_tmp = 0.7
dj_db_tmp, dj_dw_tmp = compute_gradient_logistic_reg(X_tmp, y_tmp, w_tmp, b_tmp,
, lambda_tmp)

print(f"dj_db: {dj_db_tmp}", )
print(f"Regularized dj_dw:\n {dj_dw_tmp.tolist()}", )
```

Expected Output

```
dj_db: 0.341798994972791
Regularized dj_dw:
[0.17380012933994293, 0.32007507881566943, 0.10776313396851499]
```

Rerun over-fitting example

In []:

```
plt.close("all")
display(output)
ofit = overfit_example(True)
```

In the plot above, try out regularization on the previous example. In particular:

- Categorical (logistic regression)
 - set degree to 6, lambda to 0 (no regularization), fit the data
 - now set lambda to 1 (increase regularization), fit the data, notice the difference.
- Regression (linear regression)
 - try the same procedure.

Congratulations!

You have:

- examples of cost and gradient routines with regularization added for both linear and logistic regression
- developed some intuition on how regularization can reduce over-fitting

In []: